Nexora Products and Platforms

Nexora Systems provides a portfolio of enterprise AI infrastructure products that help organizations deploy, monitor, and optimize AI systems.

Nexora ControlHub

ControlHub is the flagship AI observability platform developed by Nexora Systems. The platform provides deep visibility into the performance and behavior of large language models and AI agents deployed in enterprise environments.

ControlHub monitors several key performance indicators including:

- response latency
- token consumption
- model cost per request
- hallucination probability
- response consistency
- error rates
- retry frequency

The platform integrates with popular AI frameworks including LangChain and supports monitoring for both open source and proprietary large language models.

ControlHub also includes automated optimization capabilities that allow organizations to dynamically adjust prompts, model selection, temperature parameters, and retrieval configurations.

Nexora AgentFlow

AgentFlow is an AI orchestration platform that allows organizations to build complex AI agents capable of performing multi step reasoning tasks. These agents can interact with external APIs, internal databases, document repositories, and enterprise tools.

AgentFlow includes features such as:

- tool calling
- multi agent workflows
- memory systems
- retrieval augmented generation pipelines
- document reasoning chains

The platform allows developers to build AI systems capable of executing complex workflows such as financial analysis, legal document review, recruitment screening, and automated reporting.

Nexora VectorCloud

VectorCloud is a distributed vector database platform optimized for large scale semantic search workloads. The system is designed to store billions of embeddings while maintaining low latency query performance.

VectorCloud supports hybrid search techniques including keyword search, vector similarity search, and metadata filtering.

The platform is commonly used to power retrieval augmented generation pipelines where AI agents must search through large collections of documents to answer user queries.